

CSE 575 Homework 1

Sunayana Gupta

1222389090

- [4pts] implement your own version of the method of least-squares, compute and report θ_0 and θ_1 that minimize the residual sum of squares,

$$\sum_{i=1}^N \frac{1}{2} (y(i) - h_{\theta}(x(i)))^2$$

Solution:

```
def least_squares(X, y, m, c):
    X_mean = np.mean(X)
    Y_mean = np.mean(y)

    num = 0
    den = 0
    for i in range(len(X)):
        num += (X[i] - X_mean) * (y[i] - Y_mean)
        den += (X[i] - X_mean) ** 2
    m = num / den
    c = Y_mean - m * X_mean

    return [c, m]
```

Theta = [array([152.91886183]), array([938.23786125])]

- [4pts] implement your own version of the gradient descent algorithm, compute and report θ_0 and θ_1 that minimize the mean squared error

$$\sum_{i=1}^N \frac{1}{2} (y(i) - h\theta(x(i)))^2$$

Solution:

```
def gradient_descent(X, y, theta, alpha, n_iter):
    n_samples = len(y)
    X = np.hstack((np.ones((n_samples, 1)), X))
    y = y[:, np.newaxis]
    for i in range(n_iter):
        theta = theta - ((alpha/n_samples) * X.T @ (X @
theta - y))
    return theta
```

Theta = [[152.918864] [938.23328304]]

- [2pts] derive the analytical expression of the gradient if the loss is defined as

$$\sum_{i=1}^N \frac{1}{2} (y(i) - h_{\theta}(x(i)))^2 + \frac{\lambda}{2} \|\theta\|_2^2, \text{ where } \theta = [\theta_0, \theta_1]^T$$

Q- Derive the analytical expression of the gradient if loss is defined as

$$\sum_{i=1}^N \frac{1}{2} (y^{(i)} - h_0(x^{(i)}))^2 + \frac{\lambda}{2} \|\theta\|_2^2$$

where $\theta = [\theta_0, \theta_1]^T$, $h_0(x) = \theta_0 + \theta_1 x$

Solution- We know that,

$$\theta_{1, \text{new}} = \theta_1 - \eta \frac{dL}{d\theta_1}$$

$$\theta_{0, \text{new}} = \theta_0 - \eta \frac{dL}{d\theta_0}$$

Given that,

$$L = \sum_{i=1}^N \frac{1}{2} (\theta_1 x^{(i)} + \theta_0 - y^{(i)})^2 + \frac{\lambda}{2} (\theta_1^2 + \theta_0^2)$$

$$\frac{dL}{d\theta_1} = \frac{1}{2} \times 2 \sum_{i=1}^N (h(x^{(i)}) - y^{(i)}) x + \lambda \theta_1$$

$$\frac{dL}{d\theta_0} = \frac{1}{2} \times 2 \sum_{i=1}^N (h(x^{(i)}) - y^{(i)}) + \lambda \theta_0$$

Therefore,

$$\theta_{1, \text{new}} = \theta_1 - \eta \left(\sum_{i=1}^N (h(x^{(i)}) - y^{(i)}) x + \lambda \theta_1 \right)$$

$$\theta_{0, \text{new}} = \theta_0 - \eta \left(\sum_{i=1}^N (h(x^{(i)}) - y^{(i)}) + \lambda \theta_0 \right)$$

Complete code –

```
# -*- coding: utf-8 -*-
"""CSE575_HW1.ipynb

Automatically generated by Colaboratory.

Original file is located at
    https://colab.research.google.com/drive/1kguZUDVX60xyfhpDofzPyphXLyagMQvH

# Linear regression [10 pts]

In this homework, you will implement solution algorithms for linear regression.

## Import libraries
Let's begin by importing some libraries.
"""

# Commented out IPython magic to ensure Python compatibility.
print(__doc__)
import matplotlib.pyplot as plt
import numpy as np
from sklearn import datasets
# %matplotlib inline

"""## Load dataset

Now, we are importing a dataset of diabetes. You can check the details on this dataset
here: https://scikit-learn.org/stable/datasets/toy\_dataset.html#diabetes-dataset.

The dataset consists of 442 observations with 10 attributes ($X$) that may affect the
progression of diabetes ($y$). Ten baseline variables, age, sex, body mass index,
average blood pressure, and six blood serum measurements were obtained for each of $n$
= 442 diabetes patients, as well as the response of interest, a quantitative measure
of disease progression one year after baseline.
"""

# Load the diabetes dataset
diabetes_X, diabetes_y = datasets.load_diabetes(return_X_y=True)
print('The shape of the input features:',diabetes_X.shape)
print('The shape of the output variable:',diabetes_y.shape)

"""We will choose just one attribute from the ten attributes as an input variable."""

# Use only one feature
diabetes_X_one = diabetes_X[:, np.newaxis, 2]
print(diabetes_X_one.shape)
```

```
"""## Dataset split
```

Now, we split the dataset into two parts: training set and test set.

```
- training set: 422 samples
- test set: 20 samples
"""
```

```
# Split the data into training/testing sets
```

```
diabetes_X_train = diabetes_X_one[:-20]
```

```
diabetes_X_test = diabetes_X_one[-20:]
```

```
# Split the targets into training/testing sets
```

```
diabetes_y_train = diabetes_y[:-20]
```

```
diabetes_y_test = diabetes_y[-20:]
```

```
print('Training input variable shape:', diabetes_X_train.shape)
```

```
print('Test input variable shape:', diabetes_X_test.shape)
```

```
from matplotlib import pyplot as plt
```

```
plt.rcParams["figure.figsize"] = [7.00, 3.50]
```

```
plt.rcParams["figure.autolayout"] = True
```

```
plt.grid()
```

```
plt.plot(diabetes_X_one, diabetes_y, marker="o", markersize=5, markeredgecolor="red",
markerfacecolor="green")
```

```
plt.show()
```

```
"""## Linear regression
```

Assume that we have a hypothesis $h_{\theta}(x) = \theta_0 + \theta_1 x$.

Your tasks:

- [4pts] implement your own version of the method of least-squares, compute and report θ_0 and θ_1 that minimize the residual sum of squares,
$$\sum_{i=1}^N \frac{1}{2} (y^{(i)} - h_{\theta}(x^{(i)}))^2$$

- [4pts] implement your own version of the gradient descent algorithm, compute and report θ_0 and θ_1 that minimize the mean squared error
$$\sum_{i=1}^N \frac{1}{N} (y^{(i)} - h_{\theta}(x^{(i)}))^2$$

- [2pts] derive the analytical expression of the gradient if the loss is defined as
$$\sum_{i=1}^N \frac{1}{2} (y^{(i)} - h_{\theta}(x^{(i)}))^2 + \frac{\lambda}{2} \|\theta\|_2^2$$
,
where $\theta = [\theta_0, \theta_1]^T$

To check whether your computation is correct, consider using an API such as Scikit learn linearregression.

```
"""
```

```
def least_squares(X,y,m,c):
```

```
    X_mean = np.mean(X)
```

```
    Y_mean = np.mean(y)
```

```
    num = 0
```

```
    den = 0
```

```
    for i in range(len(X)):
```

```
        num += (X[i] - X_mean)*(y[i] - Y_mean)
```

```
        den += (X[i] - X_mean)**2
```

```
    m = num / den
```

```
    c = Y_mean - m*X_mean
```

```
    return [c, m]
```

```
def gradient_descent(X, y, theta, alpha, n_iter):
```

```
    n_samples = len(y)
```

```
    X = np.hstack((np.ones((n_samples, 1)), X))
```

```
    y = y[:, np.newaxis]
```

```
    for i in range(n_iter):
```

```
        theta = theta - ((alpha/n_samples) *X.T @ (X @ theta - y))
```

```
    return theta
```

```
def predict(X,theta):
```

```
    n_samples = len(X)
```

```
    X = np.hstack((np.ones((n_samples, 1)), X))
```

```
    y_pred = X @ theta
```

```
    return y_pred
```

```
def score(X, y,theta):
```

```
    n_samples = np.size(X, 0)
```

```
    X = np.hstack((np.ones((n_samples, 1)), X))
```

```
    y = y[:, np.newaxis]
```

```
    y_pred = X @ theta
```

```
    score = 1 - (((y - y_pred)**2).sum() / ((y - y.mean())**2).sum())
```

```
    return score
```

```
n_iters = 60000
```

```
n_features = np.size(diabetes_X_train, 1)
```

```
learning_rate = 0.09
```

```
m=0
```

```
c=0
```

```
(optimal_params) = least_squares(diabetes_X_train,diabetes_y_train , m,c)
```

```
print("Optimal parameters are: \n", optimal_params, "\n")
```

```
accuracy= score(diabetes_X_test,diabetes_y_test, optimal_params)
```

```

print("Score: ", accuracy, "\n")

n_iters = 60000
n_features = np.size(diabetes_X_train, 1)
learning_rate = 0.09
theta = np.zeros((n_features + 1, 1))
(optimal_params) = gradient_descent(diabetes_X_train,diabetes_y_train , theta,
learning_rate, n_iters)
print("Optimal parameters are: \n", optimal_params, "\n")
accuracy= score(diabetes_X_test,diabetes_y_test, optimal_params)
print("Score: ", accuracy, "\n")

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression as lr

our_parameters = gradient_descent(diabetes_X_train,diabetes_y_train , theta,
learning_rate, n_iters)
sklearn_regressor = lr().fit(diabetes_X_train, diabetes_y_train)
our_train_accuracy = score(diabetes_X_train,diabetes_y_train, optimal_params)
sklearn_train_accuracy = sklearn_regressor.score(diabetes_X_train, diabetes_y_train)
our_test_accuracy = score(diabetes_X_test,diabetes_y_test, optimal_params)
sklearn_test_accuracy = sklearn_regressor.score(diabetes_X_test, diabetes_y_test)

print("sklearn_train_accuracy: " ,sklearn_train_accuracy)
print("our_train_accuracy: ", our_train_accuracy)
print("sklearn_test_accuracy: ",sklearn_test_accuracy)
print("our_test_accuracy: " ,our_test_accuracy)

```